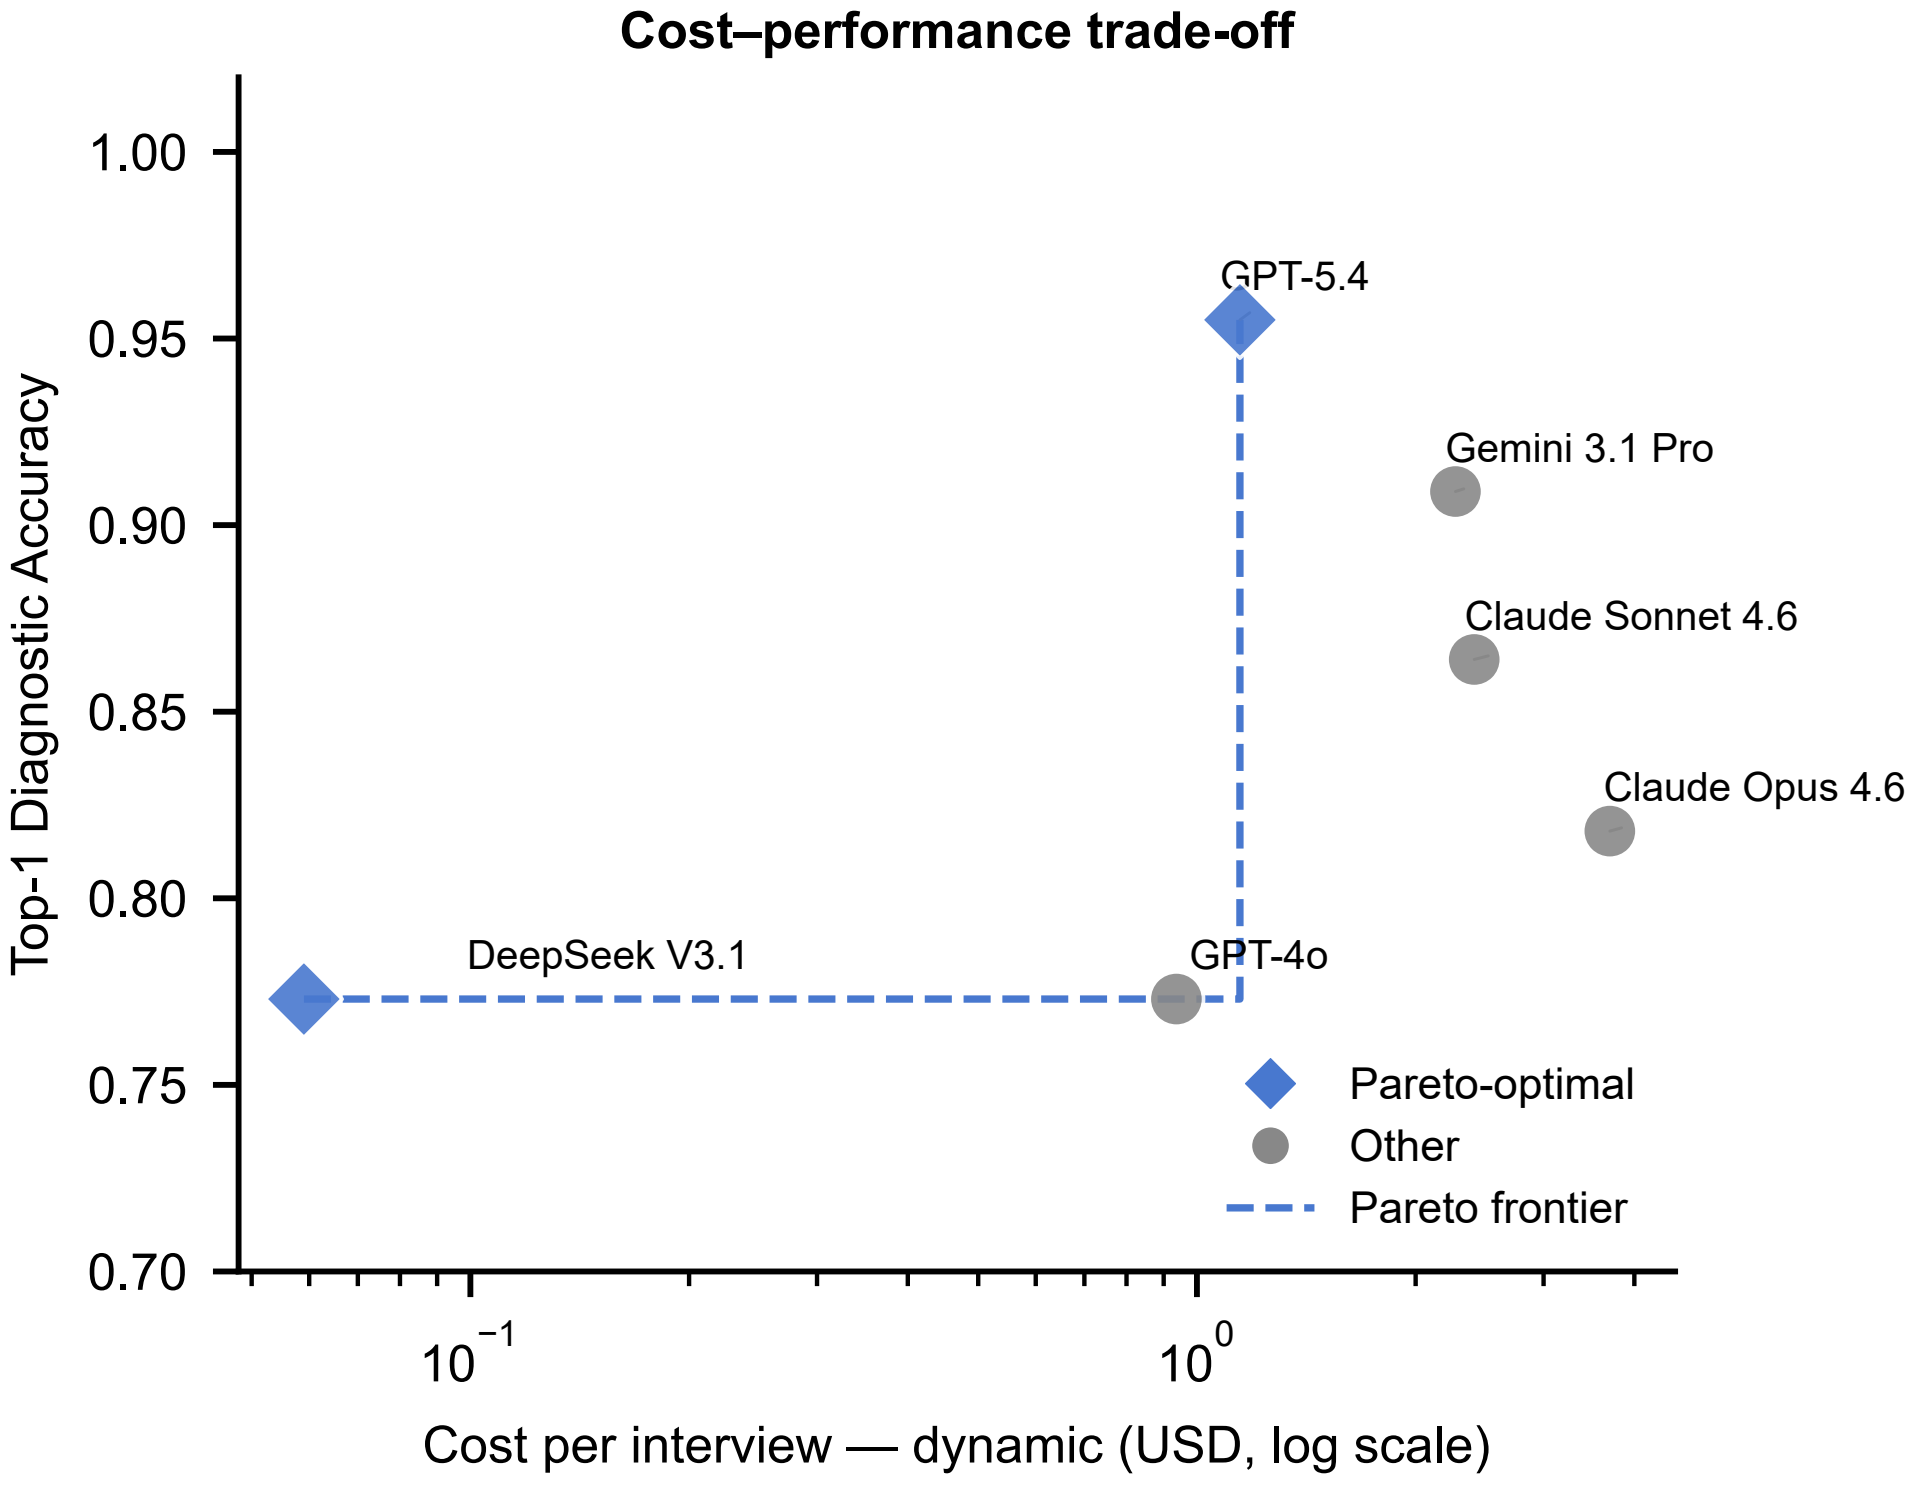
Cost–performance trade-off
1.00
0.95
0.90
0.85
0.80
0.75
0.70
Top-1 Diagnostic Accuracy
GPT-5.4
Gemini 3.1 Pro
Claude Sonnet 4.6
Claude Opus 4.6
DeepSeek V3.1
GPT-4o
Pareto-optimal
Other
Pareto frontier
10^{-1}
10^{0}
Cost per interview — dynamic (USD, log scale)